

CS342: Operating Systems Lab.
Department of Computer Science and Engineering,
Indian Institute of Technology, Guwahati
North Guwahati, Assam 781 039

Exercise T-01

OS Lessons: Busy wait, Thread states, Timer Alarms
Rating: Moderately difficult

Last update: 20 February 2017

Primary reference: Pintos by Ben Pfaff (Referred below as PintDoc)

Our focus continues to include training with the *Developmental Tools* (Appendix F) and *Debugging Tools* (Appendix E) suggested in PintDoc. To support and set a Pintos specific goal we will add a kernel augmentation exercise from PintDoc. This exercise is described in Section 2.2.2 *Alarm Clock*. We also add the first two paragraphs of Section 2.2.3 *Priority Scheduling* in this exercise. Rest of this document adds to the description provided in PintDoc to guide the students towards a solution.

We do recognize that at this time of the semester CS342 students are yet to learn concurrent programming topics; specifically, how the activities happening independently yet simultaneously coordinate their actions inside a computer – try walking in a crowded room with everyone blindfolded. This document explains how to organize your program and active entities called *threads*. You must use the tools listed in the previous paragraph to locate functions in Pintos code: understand them and use them as needed to develop the solution code.

A typical OS kernel hosts a number of threads. A thread is an independent activity capable of executing its instructions on a computer. We will ignore all issues that concern with the creation, resource allocation and termination of a thread to focus on how a thread is scheduled to run on a computer processor.

Assume a single processor computer; only one thread can be executing at any given time. The thread is said to be in state `THREAD_RUNNING`. The other threads must wait for their turns. A threads waiting for its turn to use the processor is a ready thread. And, its state is `THREAD_READY`. The threads that are not seeking use the processor at a point in time are blocked threads in state `THREAD_BLOCKED`. These threads may be waiting for an event. In the exercise for this week, the events of interest to the waiting/blocked threads are the timer events. These are also called timer alarms or just alarms.

37 To ensure that all ready threads receive a fair use of the processor time, the ready
38 threads are scheduled by a mechanism whose details we will ignore in this exercise.
39 All that matters to us is the fact that each thread gets a small amount of the
40 processor time (4 ticks of a clock) to run before being told to wait for the next turn.

41 As explained in PintDoc, time is measured inside a processor as ticks. A tick counting
42 mechanism is included in the given PintOS code. A thread that wants to sleep and
43 not do anything for a certain number of ticks, may pass or waste its turn till the time
44 has progressed adequately. The current implementation wastes the time by calling
45 `thread_yield()` – the thread lets some other unknown thread use the
46 processor time. This is bad because the scheduler must keep asking every thread
47 frequently and/or periodically if the thread wants to use the processor. Further, the
48 thread being asked to use the slack time may not be interested – the thread will in
49 turn pass the option to yet another thread. The chain of threads unwilling to use the
50 available processor time slot may be arbitrarily long. The blind `thread_yield()`
51 does not direct the request to make use of the free processor time to the threads
52 willing to use the time. An obvious waste of the processor time.

53 A better method is to make the threads waiting for alarms blocked till the required
54 number of ticks have been counted. The better algorithm is outlined as the changed
55 code of function `timer_sleep()` in file `devices/timer.c`:

```
56 /* Sleeps for approximately TICKS timer ticks.  
57    Interrupts must be turned on. */  
58 void  
59 timer_sleep (int64_t ticks)  
60 {  
61     int64_t start = timer_ticks ();  
62     int64_t wakeup_at = start + ticks;  
63  
64     ASSERT (intr_get_level () == INTR_ON);  
65  
66     /* Put the thread to sleep in timer sleep queue */  
67     thread_priority_temporarily_up ();  
68     thread_block_till (wakeup_at, before);  
69  
70     /* original code -- to be decommissioned */  
71     while (timer_elapsed (start) < ticks)  
72     thread_yield (); */  
73  
74     /* Thread must quit sleep and also free its successor  
75        if that thread needs to wakeup at the same time. */  
76     thread_set_next_wakeup ();  
77     thread_priority_restore ();  
78 }  
79
```

Outline of the proposed algorithm

A thread planning to sleep must arrange to be woken up at a suitable time. Once the arrangement is made, the thread can block itself. The OS scheduler will not consider the blocked threads for allocation of time slot for execution on the processor.

The OS, however, must keep track of the time at which a blocked thread needs unblocking. Obviously, the thread of interest for this purpose is the one that has the earliest wakeup time among those who are blocked/sleeping for their timer alarms. The OS only needs to watch one (wakeup) time: the time for the next wakeup.

The time for the next alarm that the OS tracks can change only under two conditions (and, no other!):

- A new thread joins the set of sleeping threads and has wakeup time before the current earliest/next wakeup time. This change can be made by the thread planning to sleep before it actually blocks. Function `thread_block_till (wakeup_at, before)` implements this requirement. Or,
- A sleeping thread is woken at the end of its sleep time. The thread must look at all the remaining threads that are sleeping to find the nearest time for the next wakeup. Set the system to perform next wakeup at this new time. This is among the jobs for function `thread_set_next_wakeup ()`.

We can let function `thread_tick()` do *one* unblock of *one* sleeping thread at each wakeup time. (Significance of emphasis on words *one* will become clear a little later. A simple way to impose this restriction is by changing next-wake-time after unblocking one waiting thread. This unblocked thread when prudent will correctly set the next-wakeup-time.)

There are more details to attend but we must wait for them.

If you are wondering about two other functions

`thread_priority_temporarily_up()` and `thread_priority_restore()`, the answer has many parts:

- The former function is called when a thread is about to block, so its temporary higher priority is not too restrictive to the other threads.
- Note there are several sleeping threads. So we need a list of sleeping (blocked) threads. This list is a shared list among all sleeping threads. A shared resource such as this list is only usable serially on one-thread-at-a-time basis (mutual exclusion). To keep the wall-clock time for the use-duration short we may wish to run the threads using the list at the highest priority levels.
- The thread that was blocked was good to the other threads as other threads could use the freed processor time. When a thread wakes, the benefitted

118 threads may be nice to it in-turn. We delay the restoration of the priority
119 levels to the last step in the proposed algorithm.

120 Task 1:

121 It is instructive to work on these two priority-changing functions first to gain
122 familiarity with the kernel code. Search the code using the development tools
123 (`cscope` and/or `ctags`) to see how a `ready_list` is used to record all ready
124 threads in the system. It will take about a dozen lines of new code and changes to
125 implement functions `thread_priority_temporarily_up()` and
126 `thread_priority_restore()`. You will require a new member in `struct`
127 `thread` to save priority.

128 How to Determine If Your Changes Are Right?

129 When you run `make check` to test unmodified `Pintos` implementation under
130 directory `pintos/src/thread` you have only 19 failed tests. On completion of
131 the priority changing functions (we are still using the original while-loop with
132 `thread_yield()` to wait for the specified number of ticks in function
133 `timer_sleep()`), your success count for command `make check` should remain
134 unchanged.

135 Also, note that on completion of this stage, `ready_list` is a sorted list sorted by
136 thread priority. And, you would have written a function of type `list_less_func`
137 to compare threads in `ready_list` for their correct position in the list.

138 Task 2:

139 Now is the time to introduce another sorted list to hold threads that are blocked on
140 their timer alarms. For convenience, we refer to it as `list_sleepers`. Remember
141 that this list should only be accessed by one thread at a time. Only after a thread has
142 satisfied its needs is `list_sleepers` allowed to be accessed by another thread.

143 Let us now describe an algorithm for other two functions: `thread_block_till`
144 (`wakeup_at`, `before`) and `thread_set_next_wakeup()`. Function
145 parameter `before` is a function to compare the threads in `list_sleepers` and its
146 prototype is `bool before (const struct list_elem *a, const`
147 `struct list_elem *b, void *aux UNUSED)` also aliased as
148 `list_less_func`.

149 As `list_sleepers` may be a long list, some list functions described in
150 `kernel/list.h` may take multiple ticks to complete. This may cause other
151 threads to be unblocked and they too may access `list_sleepers` before a previous
152 thread has finished its use of the list. This is not safe and we do not want to let it
153 happen.

154 We must control access to the list by a proper synchronization tool from `synch.h`.
155 The tool will deny access to the list by other threads while one thread is changing the

156 list. To keep these “busy” durations short, we have already arranged for temporarily
157 increase of the scheduling priority of the relevant threads around these two
158 functions.

159 The thread must release the synch object before it enters the block state. If access to
160 list `sleepers` is not free when the “controlling” thread is blocked, other threads
161 will remain unable to access list `sleepers`. Even a tiny time (just a few
162 instructions) between the release of a synchronization control and the actual block
163 of the thread is a potential trouble spot. You must manage it. Fortunately, the task is
164 not too difficult. Surely, there are available solutions in the kernel code itself for you
165 to learn the trick.

166 Task 3:

167 When a `sleepers` is woken (unblocked), the situation is trickier. On any given
168 single timer tick, there may be several sleeping threads seeking to wake up. How do
169 you attend to all of them within a single tick by calling `thread_unblock()` for
170 each of them?

171 We suggest, you write code for this in function `thread_set_next_wakeup ()`
172 which was primarily designed to remove thread `thread_current()` from list
173 `sleepers` when the unblocked thread is scheduled after it is woken/unblocked
174 from its sleep by function `thread_ticks()`.

175 This is not all. The function must also make sure that if there are other sleeping
176 threads that have the same wakeup time as the thread just woken up, then they are
177 a not left blocked. If a sleeping thread with the matching wakeup time is left waiting,
178 then the waiting thread can only wakeup on a later timer tick. On the other hand, we
179 cannot spend too much time waking all such threads because we do not want to
180 hold the timer-tick in disabled state for too long. A missed tick would drift the
181 system clock. We solved the problem by limiting to one thread unblock in function
182 `thread_ticks()`.

183 The task of unblocking other simultaneous wakeups is easily shared among the
184 multiple threads. Each unblocked thread is required to unblock one thread that has
185 the same wakeup time as itself. This action will recursively unblock all threads which
186 have the matching wakeup times. All except the first call to `thread_unblock()`
187 will be from function `thread_set_next_wakeup ()`. An unblocked thread
188 that does not unblock another thread sets the next-wake-up time. Suffice to remind
189 that this time was set far into the future to prevent multiple unblocks from function
190 `thread_tick()`.

191 Never unblock a sleeping thread if there is another unblocked thread with an earlier
192 wakeup time.

193 **Reminders**
194 Be sure to check out and check-in the versions of your program from and into your
195 version control repository. Promptly and properly annotate your program with
196 comments.

197 You may also note that we have implemented basic thread priority idea as it makes
198 this implementation simpler. This is a bit more than PintDoc 2.2.2 *Alarm Clock* target.
199 You may wish to see PintOS-T03 guide for an alternate implementation guide.

200 **Results of make check Command Execution:**

```
201 pass tests/threads/alarm-single
202 pass tests/threads/alarm-multiple
203 pass tests/threads/alarm-simultaneous
204 pass tests/threads/alarm-priority
205 pass tests/threads/alarm-zero
206 pass tests/threads/alarm-negative
207 FAIL tests/threads/priority-change
208 FAIL tests/threads/priority-donate-one
209 FAIL tests/threads/priority-donate-multiple
210 FAIL tests/threads/priority-donate-multiple2
211 FAIL tests/threads/priority-donate-nest
212 FAIL tests/threads/priority-donate-sema
213 FAIL tests/threads/priority-donate-lower
214 FAIL tests/threads/priority-fifo
215 FAIL tests/threads/priority-preempt
216 FAIL tests/threads/priority-sema
217 FAIL tests/threads/priority-condvar
218 FAIL tests/threads/priority-donate-chain
219 FAIL tests/threads/mlfqs-load-1
220 FAIL tests/threads/mlfqs-load-60
221 FAIL tests/threads/mlfqs-load-avg
222 FAIL tests/threads/mlfqs-recent-1
223 pass tests/threads/mlfqs-fair-2
224 pass tests/threads/mlfqs-fair-20
225 FAIL tests/threads/mlfqs-nice-2
226 FAIL tests/threads/mlfqs-nice-10
227 FAIL tests/threads/mlfqs-block
228 19 of 27 tests failed.
229 make[1]: *** [check] Error 1
230
```

231 **Contributing Authors:**

232 Vishv Malhotra, Gautam Barua, Rashmi Dutta Baruah